

# Distributed Systems and Algorithms

Martin Quinson <martin.quinson@loria.fr>  
Abdelkader Lahmadi <abdelkader.lahmadi@loria.fr>  
Olivier Festor <olivier.festor@inria.fr>

LORIA – INRIA Nancy Grand Est

2018-2019

(compiled on: September 26, 2019)

# Introduction

## Course Goals

- ▶ Introduce existing distributed systems, from a theoretical point of view
  - ▶ Basic concepts
  - ▶ Main issues, problems and solutions

## Motivations

- ▶ Distributed Systems more and more mainstream
- ▶ Interesting algorithmic issues

# Administrativae

## Contents

- ▶ Quick recap of distributed algorithmic and Internet
- ▶ Present several innovative distributed systems

## Evaluation: test on desk (*partiel*)





- ▶ **What:** quiz about the lectures
  - ▶ *Know* the algorithms introduced in lectures
  - ▶ Be able to *recognize* principle of classical algorithm designs
  - ▶ Be able to *discuss* the validity of an approach to a problem

## References: Courses on Internet

- ▶ **Algorithmique et techniques de base des systèmes répartis** (S. Krakowiak)  
Foundations of distributed systems (in French).  
<http://sardes.inrialpes.fr/~krakowia/Enseignement/M2R-SL/SR/>
- ▶ **Distributed Systems** (Shenker, Stoica; University of California, Berkley)  
A bit of everything, emphasis on Brewer's conjecture.  
<http://inst.eecs.berkeley.edu/~cs194>
- ▶ **Peer-to-Peer Networks** (Jussi Kangasharju)  
Peer-to-peer systems.  
<http://www.cs.helsinki.fi/u/jakangas/Teaching/p2p-08f.html>
- ▶ **Advanced Operating Systems** (Neeraj Mittal)  
Very good presentation of the theoretical foundations.  
<http://www.utdallas.edu/~neerajm/cs6378f09>
- ▶ **Grid Computing WS 09/10** (E. Jessen, M. Gerndt)  
Grid and Cloud computing.  
<http://www.lrr.in.tum.de/~gerndt/home/Teaching/WS2009/GridComputing/GridComputing.htm>

# References: Books

- ▶ Coulouris, Jean et Kindberg. **Distributed Systems: Concepts and Design**.
- ▶ Tannenbaum, Steen. **Distributed Systems: Principles and Paradigms**.
- ▶ V. K. Garg. **Elements of Distributed Computing**.
- ▶ Ralf Steinmetz, Klaus Wehrle (Eds): **Peer-to-Peer Systems and Applications**.  
<http://www.peer-to-peer.info/>



# Table of Contents

- 1 Introduction to Distributed Systems
  - ▶ What is it? Examples.
- 2 Distributed Algorithmic
  - ▶ Time and state; Ordering events; Abstract Clocks; Classical Algorithms.
- 3 Mutual Exclusion and synchronisation
  - ▶ Introduction; classical algorithms and solutions to mutual exclusion problems.
- 4 Blockchains
  - ▶ Introduction; structure and components; consensus algorithms

# Chapter 1

## Introduction

- What is a Distributed System?
- Example of Distributed Systems
- Limit between Computers and Distributed Systems

# What is a distributed system?

## Definition

*A distributed system is a collection of independent computers that appear to the users of the system as a single computer.*

— A. Tanenbaum.

↪ Set of **elements** (CPU, storage) **interconnected** by the network



- ▶ The set is more than the sum of its parts (elements do collaborate)
- ▶ Intuitive examples not from CS
  - ▶ Ant nest
  - ▶ Driving rules (cars share the road)



# What is a distributed system?

## Definition (optimistic)

*A distributed system is a collection of independent computers that appear to the users of the system as a single computer.*

— A. Tanenbaum.

- ↪ Set of **elements** (CPU, storage) **interconnected** by the network



- ▶ The set is more than the sum of its parts (elements do collaborate)
- ▶ Intuitive examples not from CS
  - ▶ Ant nest
  - ▶ Driving rules (cars share the road)

## Definition (pessimistic)

*You know you have one when the crash of a computer you never heard of stops you from getting any work done.*

— L. Lamport.

- ▶ **Interdependent** behavior of elements
  - ▶ That's not that easy
  - ▶ Failures do happen and must be dealt with

# Why would you distribute your computer system??

**Application needs:** you sometimes have to

- ▶ Collaborative work (between human beings, between corporate facilities)
- ▶ Distributed electronic devices  $\Rightarrow$  *Ubiquitous Computing and SensorNets*
- ▶ Application integration (multi-physics simulation)  $\Rightarrow$  *Grid Computing*

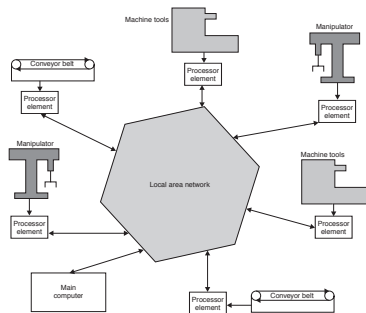
**Technical possibility creates the need**

- ▶ Cost effectiveness
  - ▶ A set of PC is less expensive than a big mainframe  $\Rightarrow$  *Cluster Computing*
  - ▶ Scale savings of mesocenter (wrt than several clusters)  $\Rightarrow$  *Cloud Computing*
- ▶ Generalized interconnections (TV, Internet, phone are converging)
  - ▶ Share storage resources  $\Rightarrow$  *Peer-to-Peer systems*
  - ▶ Share (otherwise unused) computational resources  $\Rightarrow$  *Volunteer Computing*

# Example of Distributed Systems (1/3)

## Manufacturing distributed system

- ▶ Embedded command and control systems
- ▶ Data acquisition processes



## Aircraft


- ▶ Embedded distributed systems
- ▶ Integrated Modular Avionics : processing modules interconnected via ARINC 629 bus



# Example of Distributed Systems (2/3)

## The Internet: the network of networks


- ▶ Enormous (open ended)
- ▶ No single authority  
(mapping internet is a research agenda)
- ▶ Data, audio, video; Requests, push, streams.



CoDoKi, Fig. 1.1

## Intranets

- ▶ A single authority
- ▶ Protected access  
(firewall, encrypted channels, total isolation)
- ▶ May be worldwide




CoDoKi, Fig. 1.2

# Example of Distributed Systems (3/3)


## Mobile and Ubiquitous Computing

- ▶ Portable devices
  - ▶ laptops, notebook
  - ▶ handheld, wearable devices
  - ▶ devices embedded in appliances
- ▶ Mobile computing
- ▶ Connected to Internet through fixed infrastructure



## Mobile Ad-hoc Networks (Manets)

- ▶ No fix infrastructure
  - ▶ wireless communication
  - ▶ multi-hop networking
  - ▶ long, non deterministic delays
  - ~ nodes part of infrastructure
- ▶ Nodes come and go




# Limit between Computers and Distributed Systems

## Why is this limit blurred?


- ▶ **Motivation:** endless need for power (modeling/game realism, server scalability)
- ▶ **Past solution:** Increase clock speed, more electronic gates  
(but reaching physical limits + speed linear vs. energy quadratic)
- ▶ **Current trend:** Multi-many (Multiply cores, processors and machines)

## Multi-processors systems


### Shared Memory Processor (SMP)



### Cluster System



### Distributed Systems



- ▶ SMP communicate through shared memory
- ▶ Clusters and DS communicate through classical network

# Distributed, Parallel or Concurrent??

Distributed Algorithm:  $\frac{\text{computation time}}{\text{communication time}} \rightsquigarrow 0$

- ▶ Computation negligible wrt to communications
- ▶ **Classical metric:** amount of messages (as a function of amount of nodes)

Parallel Algorithm:  $\frac{\text{computation time}}{\text{communication time}} \approx 1$

- ▶ Computation and Communication comparable
- ▶ **Classical metric:** makespan (time to completion of last processor)

Concurrent Algorithm:  $\frac{\text{computation time}}{\text{communication time}} \rightsquigarrow \infty$

- ▶ Communication negligible wrt computation ( $\text{comm time} = 0 \Rightarrow$ , multi-threading)
- ▶ **Classical metric:** speedup (how faster when using N cpus)

# Distributed, Parallel or Concurrent??

Distributed Algorithm:  $\frac{\text{computation time}}{\text{communication time}} \rightsquigarrow 0$

- ▶ Computation negligible wrt to communications
- ▶ **Classical metric**: amount of messages (as a function of amount of nodes)

Parallel Algorithm:  $\frac{\text{computation time}}{\text{communication time}} \approx 1$

- ▶ Computation and Communication comparable
- ▶ **Classical metric**: makespan (time to completion of last processor)

Concurrent Algorithm:  $\frac{\text{computation time}}{\text{communication time}} \rightsquigarrow \infty$

- ▶ Communication negligible wrt computation ( $\text{comm time} = 0 \Rightarrow$ , multi-threading)
- ▶ **Classical metric**: speedup (how faster when using N cpus)

Focus of this course: **distributed systems** (some content applies to others)



# Distributed, Parallel or Concurrent??

Distributed Algorithm:  $\frac{\text{computation time}}{\text{communication time}} \rightsquigarrow 0$

- ▶ Computation negligible wrt to communications
- ▶ Classical metric: amount of messages (as a function of amount of nodes)
- ▶ Current research agenda: P2P, consistency (distributed DB)

Parallel Algorithm:  $\frac{\text{computation time}}{\text{communication time}} \approx 1$

- ▶ Computation and Communication comparable
- ▶ Classical metric: makespan (time to completion of last processor)
- ▶ Current research agenda: Cluster & Grid & Cloud Computing, interoperability

Concurrent Algorithm:  $\frac{\text{computation time}}{\text{communication time}} \rightsquigarrow \infty$

- ▶ Communication negligible wrt computation ( $\text{comm time} = 0 \Rightarrow$ , multi-threading)
- ▶ Classical metric: speedup (how faster when using N cpus)
- ▶ Current research agenda: Lock-free, wait-free, correctness (model-checking)

Focus of this course: distributed systems (some content applies to others)

- ▶ Each domain constitutes a huge research area
- ▶ Current trend: intermixing, but strong historical heritage

# What to expect from a distributed system?

## Expected characteristics

- ▶ **Scalability:** deal with large amount of work
- ▶ **Failure tolerance:**
  - ▶ Deal with the failure of elements
  - ▶ Deal with message loss, or element performance degradation
- ▶ **Security:** Deal with malicious users (Privacy, Integrity, Deny-of-Services)
- ▶ **Adaptability:** deal with environment changes

## Expected difficulties

- ▶ **Absence of Global Clock:** there is no common notion of time
- ▶ **Absence of Shared Memory:** no process has up-to-date global knowledge
- ▶ **Failures (fail-stop or malicious):** that *will* happen
- ▶ **Delays (asynchronous):** harder to detect failures
- ▶ **Dynamism:** global knowledge even harder to get
- ▶ Human brain is (somehow) sequential. Thinking distributed is harder.